

Final Project Submission

Please fill out:

Student name:

Alvin Kimathi

- Student pace: Hybrid Full-Time
- Scheduled project review date/time: 7th June 2024
- Instructor name: Maryann Mwikali
- Blog post URL: <https://github.com/Kimathi26/Dsc-Phase-3-Project>

BUSINESS PROBLEM

SyriaTel, a tele-communications company in the Middle East has brought to our attention(Anonymous Analysts) the upward rate in which its clients are stopping usage of their product. The are looking to get ahead of the customer churn as it is easier to retain and get a return client than it is to onboard a new client.The company therefore wants to take pro-active measures to identify those likely to churn and retain them while improving customer satisfaction and reducing revenue loss.

Objectives

- Identify Key factors leading to churn
- Develop a predictive model that accurately predicts whether a customer will churn
- After visualizations propose necessary interventions

1: DATA UNDERSTANDING

We obtained our data set from Kaggle, an online hosting service which provided us with the necessary dataset to analyze this field.

1.1: Import Necessary Libraries

```
In [ ]: # Import the necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

#Preprocessing and model building

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.compose import ColumnTransformer
```

```
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

#For Models

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

#For Evaluation

from sklearn.metrics import roc_auc_score, confusion_matrix, classification_report, acc
```

1.2: Load our Dataset

```
In [ ]: data = pd.read_csv('bigml_59c28831336c6604c800002a.csv')
data.head()
```

Out []:

	state	account length	area code	phone number	international plan	voice mail plan	number vmessages	total day minutes	total day calls	total day charge	...	total eve calls	tc
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88	5
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122	12

5 rows × 21 columns

1.3: Explore and Understand our Data set

```
In [ ]: print(data.info())
print(data.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmessages                     3333 non-null   int64
```

```

7 total day minutes      3333 non-null float64
8 total day calls        3333 non-null int64
9 total day charge       3333 non-null float64
10 total eve minutes     3333 non-null float64
11 total eve calls       3333 non-null int64
12 total eve charge      3333 non-null float64
13 total night minutes   3333 non-null float64
14 total night calls     3333 non-null int64
15 total night charge    3333 non-null float64
16 total intl minutes    3333 non-null float64
17 total intl calls      3333 non-null int64
18 total intl charge     3333 non-null float64
19 customer service calls 3333 non-null int64
20 churn                 3333 non-null bool

```

dtypes: bool(1), float64(8), int64(8), object(4)

memory usage: 524.2+ KB

None

	account length	area code	number vmail messages	total day minutes	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	101.064806	437.182418	8.099010	179.775098	
std	39.822106	42.371290	13.688365	54.467389	
min	1.000000	408.000000	0.000000	0.000000	
25%	74.000000	408.000000	0.000000	143.700000	
50%	101.000000	415.000000	0.000000	179.400000	
75%	127.000000	510.000000	20.000000	216.400000	
max	243.000000	510.000000	51.000000	350.800000	

	total day calls	total day charge	total eve minutes	total eve calls	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	100.435644	30.562307	200.980348	100.114311	
std	20.069084	9.259435	50.713844	19.922625	
min	0.000000	0.000000	0.000000	0.000000	
25%	87.000000	24.430000	166.600000	87.000000	
50%	101.000000	30.500000	201.400000	100.000000	
75%	114.000000	36.790000	235.300000	114.000000	
max	165.000000	59.640000	363.700000	170.000000	

	total eve charge	total night minutes	total night calls	\
count	3333.000000	3333.000000	3333.000000	
mean	17.083540	200.872037	100.107711	
std	4.310668	50.573847	19.568609	
min	0.000000	23.200000	33.000000	
25%	14.160000	167.000000	87.000000	
50%	17.120000	201.200000	100.000000	
75%	20.000000	235.300000	113.000000	
max	30.910000	395.000000	175.000000	

	total night charge	total intl minutes	total intl calls	\
count	3333.000000	3333.000000	3333.000000	
mean	9.039325	10.237294	4.479448	
std	2.275873	2.791840	2.461214	
min	1.040000	0.000000	0.000000	
25%	7.520000	8.500000	3.000000	
50%	9.050000	10.300000	4.000000	
75%	10.590000	12.100000	6.000000	
max	17.770000	20.000000	20.000000	

	total intl charge	customer service calls
count	3333.000000	3333.000000
mean	2.764581	1.562856
std	0.753773	1.315491
min	0.000000	0.000000
25%	2.300000	1.000000
50%	2.780000	1.000000

75%	3.270000	2.000000
max	5.400000	9.000000

2: DATA CLEANING AND PREPROCESSING

2.1: Handle any Missing Values


```
In [ ]: #Check for any missing values
print("\nMissing Values in Data:")
print(data.isnull().sum())
```

```
Missing Values in Data:
state                0
account length      0
area code           0
phone number        0
international plan   0
voice mail plan     0
number vmail messages 0
total day minutes   0
total day calls      0
total day charge     0
total eve minutes   0
total eve calls      0
total eve charge     0
total night minutes 0
total night calls    0
total night charge   0
total intl minutes   0
total intl calls     0
total intl charge    0
customer service calls 0
churn                0
dtype: int64
```

3: EXPLORATIVE DATA ANALYSIS

3.1: Distribution of Churn

```
In [ ]: plt.figure(figsize=(6, 4))
sns.countplot(x='churn', data=data)
plt.title('Distribution of Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.show()
#Percentage of True/False in the distribution
churn_percentage = data['churn'].value_counts(normalize=True)* 100
churn_percentage
```




```
Out[ ]: False    85.508551  
        True     14.491449  
        Name: churn, dtype: float64
```

From this, we note that approximately 85% of the customers are still using the service while 14% have actively stopped using the service.

3.2: Univariate Analysis

3.2.1: Distribution of Numerical features

```
In [ ]: #Histograms for the numeric variables excluding boolean columns  
#Define numeric features  
numeric_features = data.select_dtypes(include=['float64', 'int64']).columns  
data[numeric_features].hist(bins=30, figsize=(20, 15))  
plt.show()
```


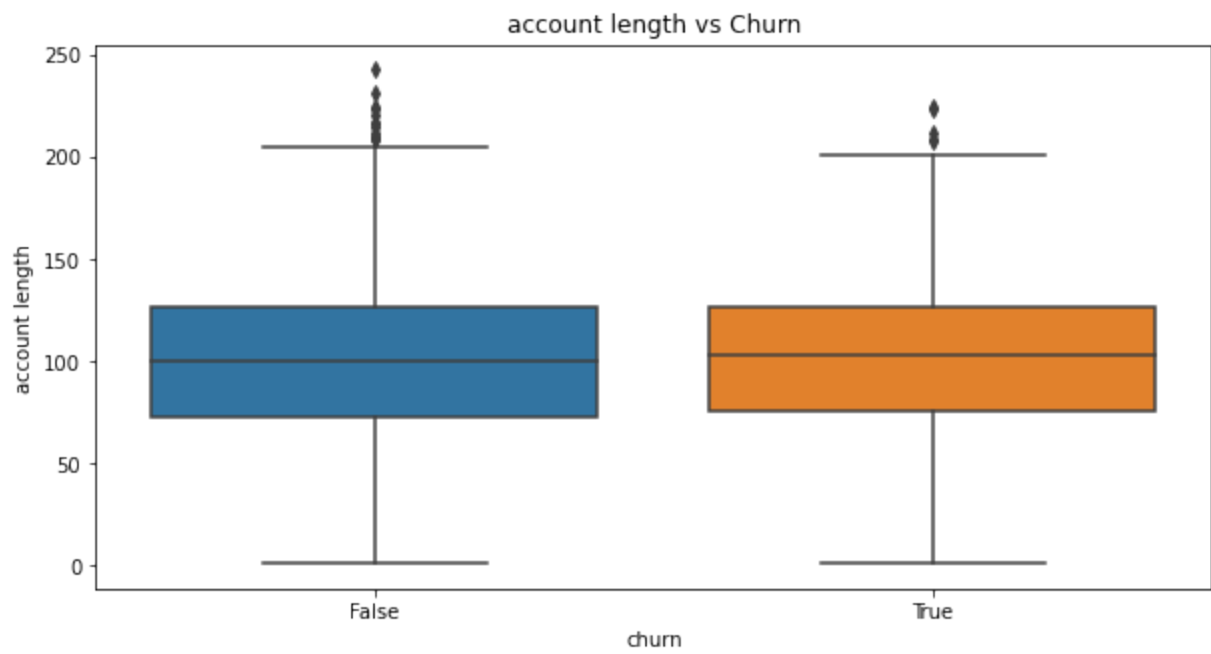


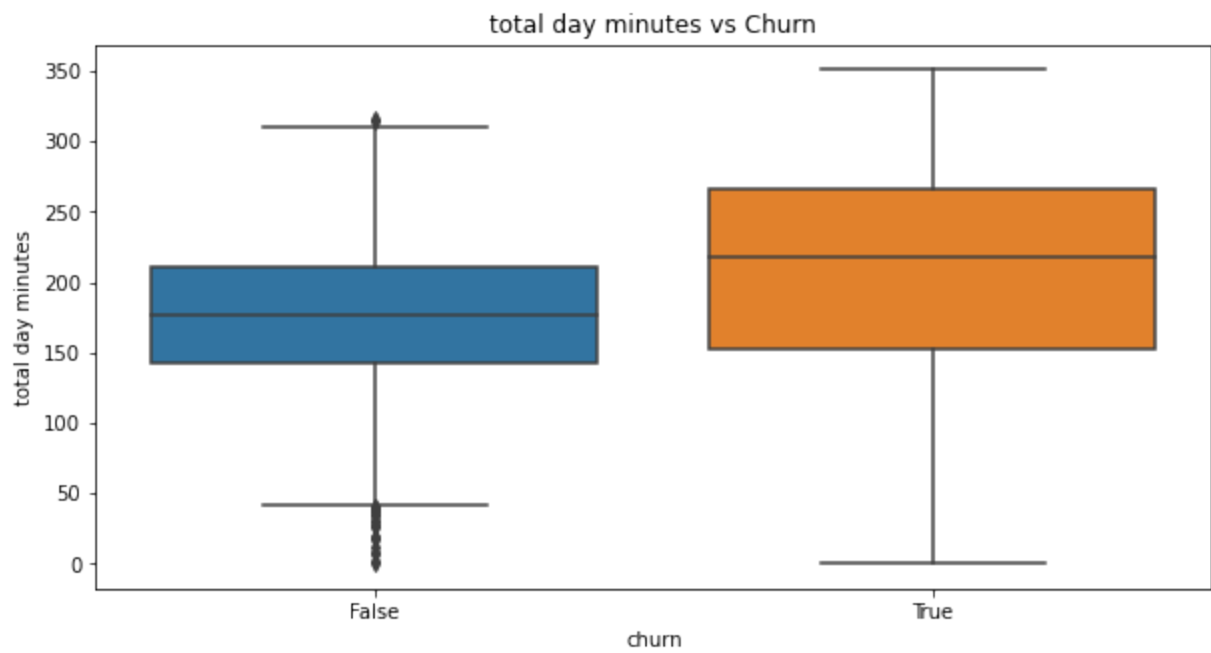
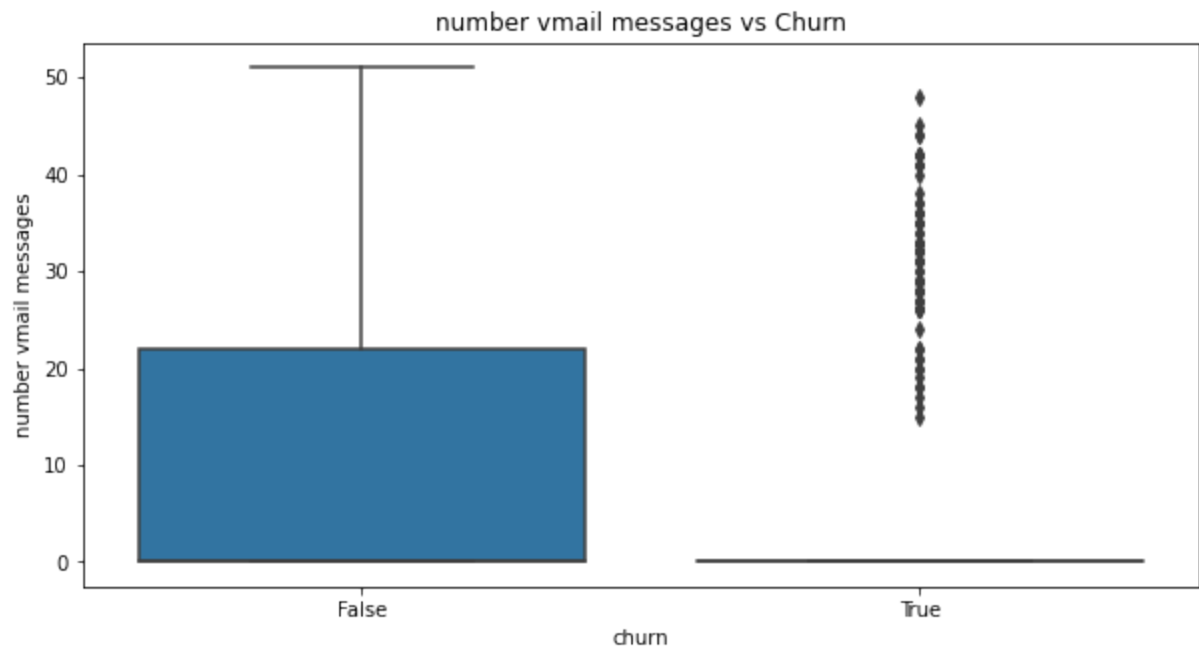
3.3: BiVariate Analysis



3.3.1: Box Plots for Numeric Features vs Churn



```
In [ ]: #exclude the 'churn' column
numeric_features = [feature for feature in numeric_features if feature != 'churn']



# Plot box plots for numerical features against Churn
for feature in numeric_features:
    plt.figure(figsize=(10, 5))
    sns.boxplot(x='churn', y=feature, data=data)
    plt.title(f'{feature} vs Churn')
    plt.show()
```





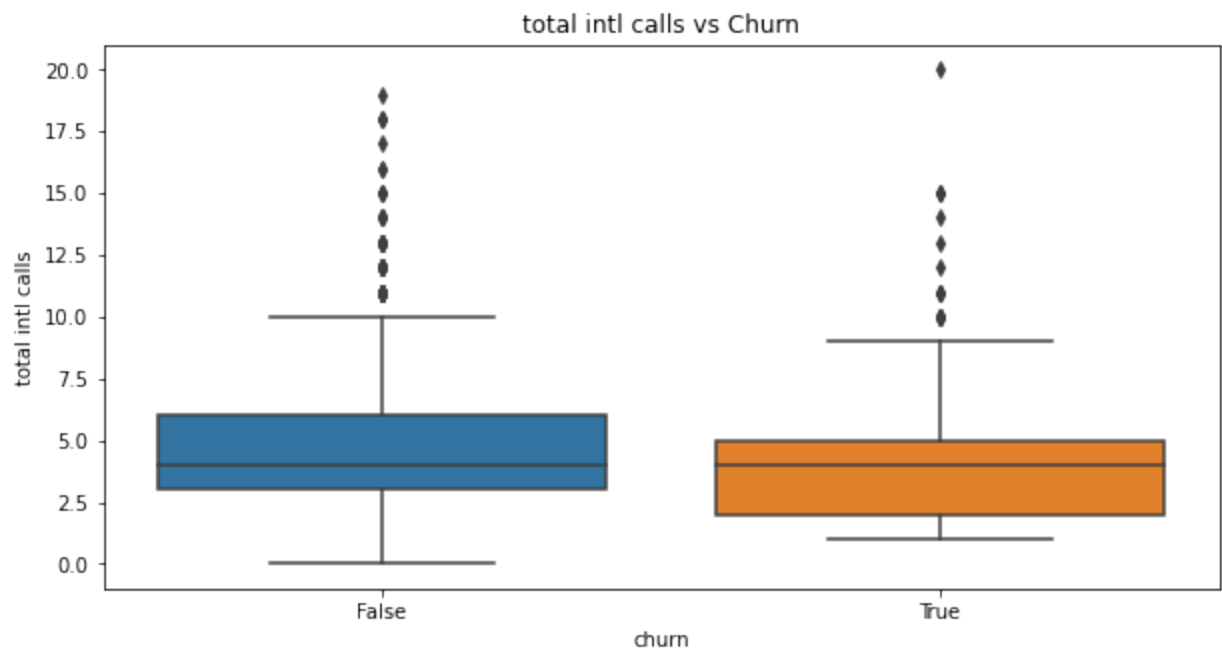
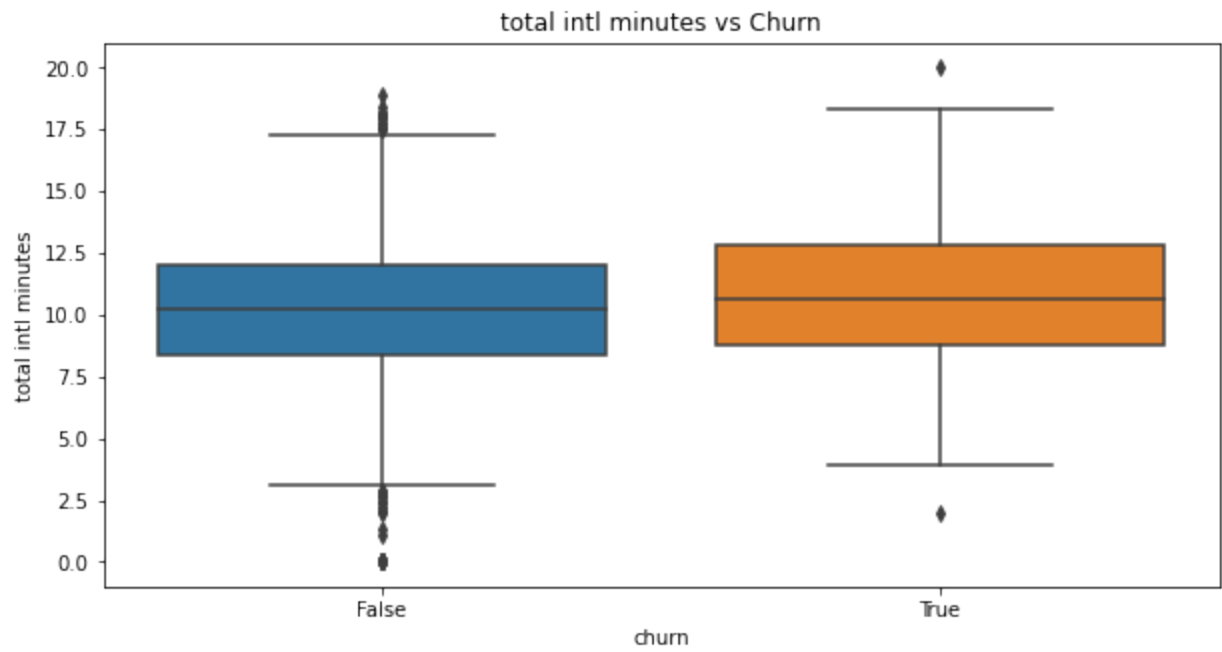


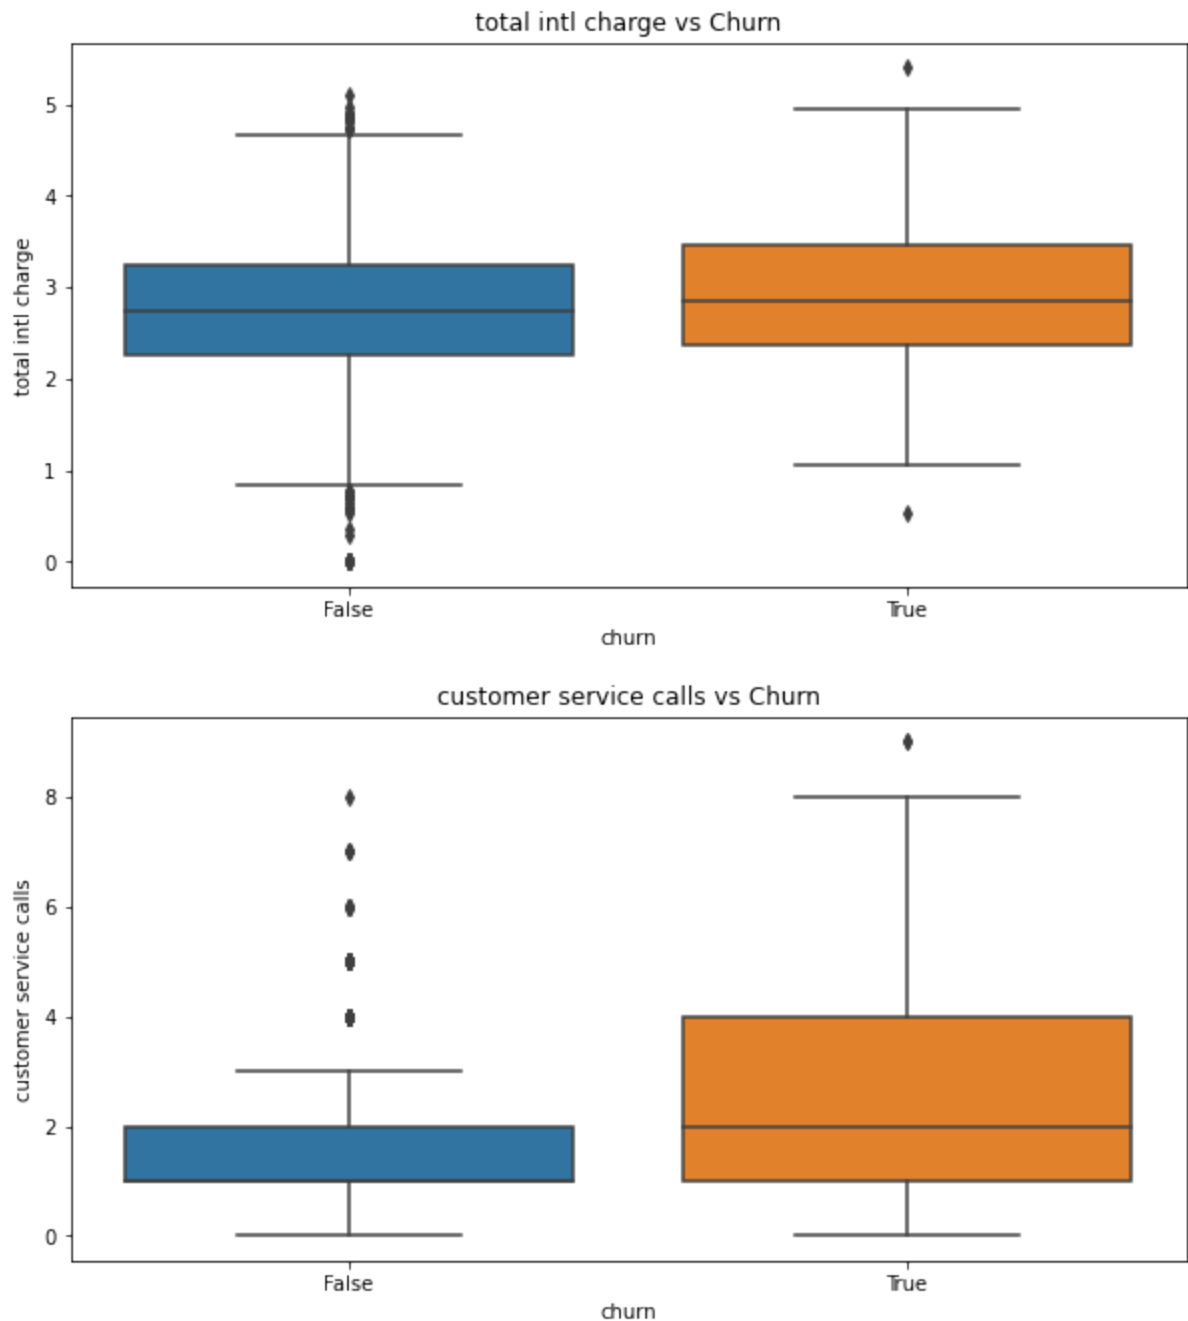















3.3.2: Count plots of Categorical features to Churn

```
In [ ]: # Identify categorical features and exclude the 'phone number' column
categorical_features = data.select_dtypes(include=['object']).columns
categorical_features = [feature for feature in categorical_features if feature != 'phone number']

# Plot cross-tabulations for categorical features against the target variable
for feature in categorical_features:
    cross_tab = pd.crosstab(data[feature], data['churn'], normalize='index')
    cross_tab.plot(kind='bar', stacked=True, figsize=(10, 5))
    plt.title(f'{feature} vs Churn')
    plt.ylabel('Proportion')
    plt.xlabel(feature)
    plt.legend(title='Churn')
    plt.show()
```








From this, one outstanding observation is that customers with an international plan are more likely to churn than those without.....probably due to higher costs if the plan is not enough to justify the cost and maybe for a short-term need as well.



3.3.3: Relationship between Numerical features and Churn



```
In [ ]: # Box plots for numeric features against the target variable(Churn)


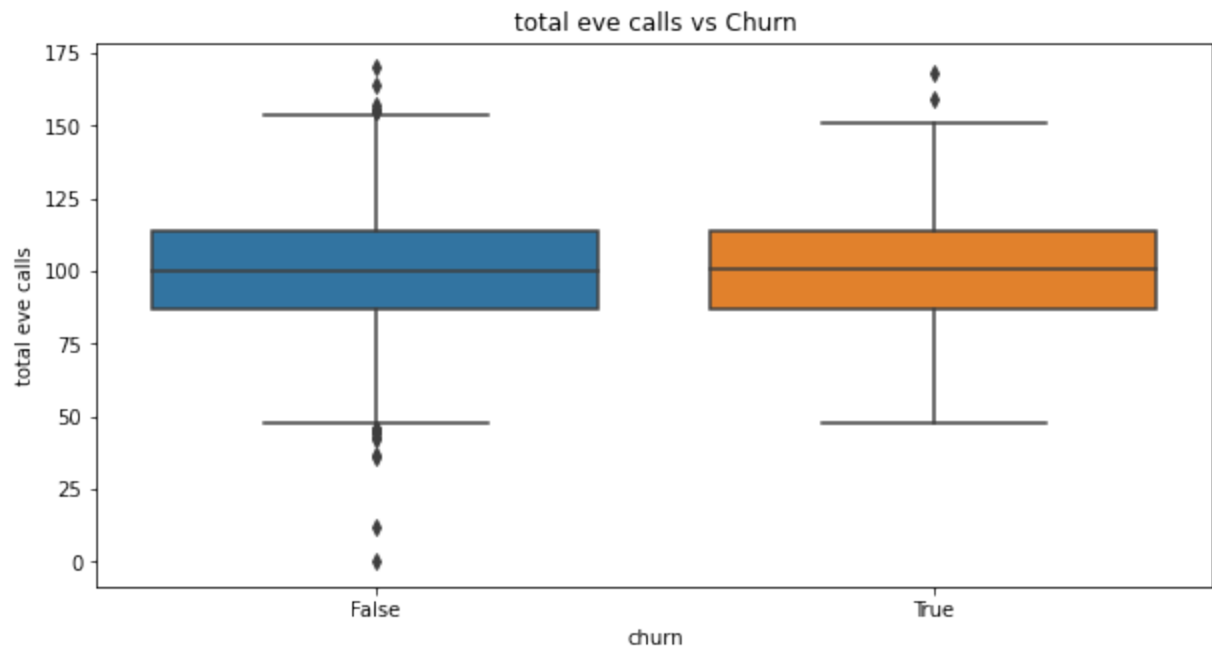
for feature in numeric_features:
    plt.figure(figsize=(10, 5))
    sns.boxplot(x='churn', y=feature, data=data)
    plt.title(f'{feature} vs Churn')
    plt.show()
```

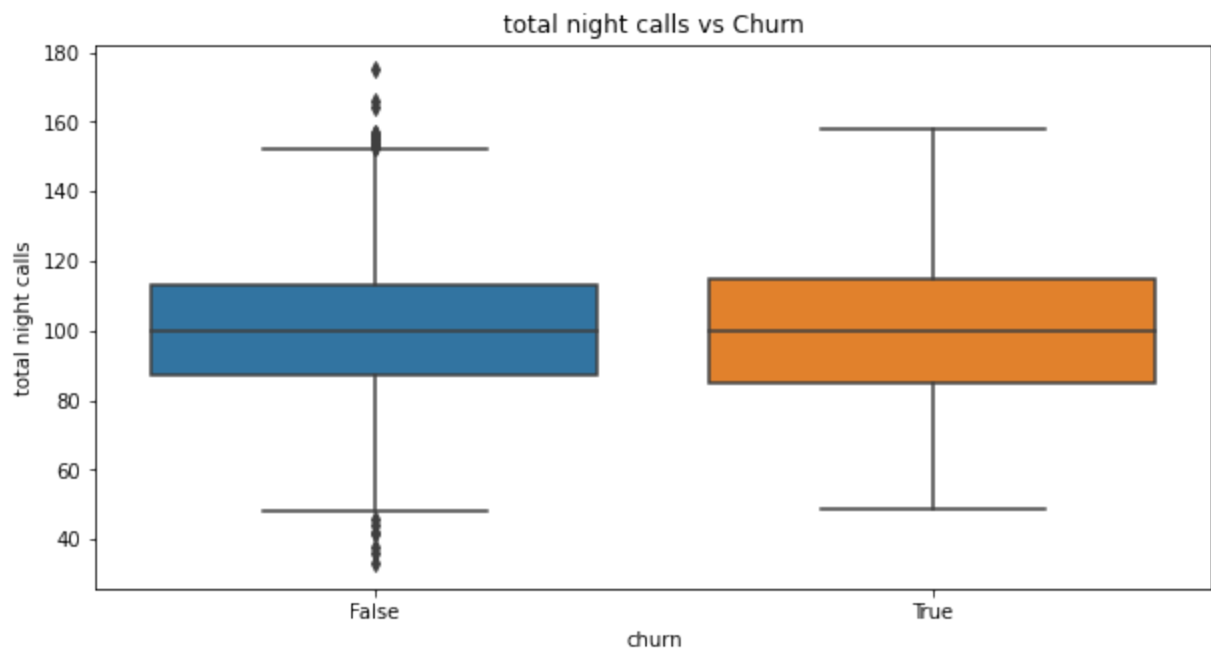

























3.4: MultiVariate Analysis

3.4.1: Correlation Matrix

```
In [ ]: # Correlation matrix
correlation_matrix = data.corr()
plt.figure(figsize=(15, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```

From this correlation heatmap we can deduce the following:

- Positive correlation between customer service calls and churn(0.21) means customers who call the service are slightly more likely to churn
- Positive correlation between total day minutes and churn(0.21) means customers who talk more during the day are also slightly likely to churn
- Positive correlation between total day charge and churn(0.21) means rates from the service provider might slightly influence the customer to churn as they continue usage. Let us also remember that correlation does not entirely equal causation as there might be other factors at play.

3.5: Encoding Categorical variables

```
In [ ]: # Identify binary categorical features for Label Encoding
binary_features = ['voice mail plan', 'international plan']

# Identify multi-class categorical features for One-Hot Encoding
multi_class_features = [feature for feature in categorical_features if feature not in binary_features]

# Apply Label Encoding to binary features
label_encoders = {}
for feature in binary_features:
    le = LabelEncoder()
    data[feature] = le.fit_transform(data[feature])
    label_encoders[feature] = le
```

```
# Apply One-Hot Encoding to multi-class features
data = pd.get_dummies(data, columns=multi_class_features, drop_first=True)

# Verify the transformations
print(data.head())
```

```
   account length  area code phone number  international plan \
0         128      415      382-4657          0
1         107      415      371-7191          0
2         137      415      358-1921          0
3          84      408      375-9999          1
4          75      415      330-6626          1

   voice mail plan  number vmail messages  total day minutes  total day calls \
0                1                25          265.1          110
1                1                26          161.6          123
2                0                 0          243.4          114
3                0                 0          299.4           71
4                0                 0          166.7          113

   total day charge  total eve minutes  ...  state_SD  state_TN  state_TX \
0         45.07          197.4  ...      0         0         0
1         27.47          195.5  ...      0         0         0
2         41.38          121.2  ...      0         0         0
3         50.90           61.9  ...      0         0         0
4         28.34          148.3  ...      0         0         0

   state_UT  state_VA  state_VT  state_WA  state_WI  state_WV  state_WY
0         0         0         0         0         0         0         0
1         0         0         0         0         0         0         0
2         0         0         0         0         0         0         0
3         0         0         0         0         0         0         0
4         0         0         0         0         0         0         0
```

[5 rows x 70 columns]

3.6: Train-Test Split Test

Train our data from the data set to be used.

```
In [ ]: # Define features and target
X = data.drop(columns=['churn', 'phone number'])
y = data['churn']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
```

3.7: Scaling our Data

To bring out improved metrics in our model building and prevent leakage as well

```
In [ ]: # Initialize the scaler
scaler = StandardScaler()

# Fit and transform the training data
X_train_scaled = scaler.fit_transform(X_train)

# Transform the test data
X_test_scaled = scaler.transform(X_test)
```

4: Model Building

4.1: Model 1 (Baseline Model-Logistic Regression)

```
In [ ]: Model_1 = LogisticRegression(max_iter=1000)
Model_1.fit(X_train_scaled, y_train)

y_pred = Model_1.predict(X_test_scaled)
y_pred_prob = Model_1.predict_proba(X_test_scaled)[:, 1]

In [ ]: print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
print("Logistic Regression Classification Report:\n", classification_report(y_test, y_p
print("Logistic Regression ROC AUC Score:", roc_auc_score(y_test, y_pred_prob))


# Plot ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
plt.figure(figsize=(10, 5))
plt.plot(fpr, tpr, label='Logistic Regression (AUC = %0.2f)' % roc_auc_score(y_test, y_
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

Logistic Regression Accuracy: 0.8575712143928036

Logistic Regression Classification Report:

	precision	recall	f1-score	support
False	0.87	0.97	0.92	566
True	0.58	0.21	0.31	101
accuracy			0.86	667
macro avg	0.73	0.59	0.61	667
weighted avg	0.83	0.86	0.83	667

Logistic Regression ROC AUC Score: 0.8164293461148235



4.2: Model 2 (Random Forest Regression Model)

```
In [ ]: # Train the random forest model
Model_2 = RandomForestClassifier(n_estimators=100, random_state=42)
Model_2.fit(X_train_scaled, y_train)

# Make predictions
y_pred_rf = Model_2.predict(X_test_scaled)
y_pred_prob_rf = Model_2.predict_proba(X_test_scaled)[: , 1]

# Evaluate the model
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Random Forest Classification Report:\n", classification_report(y_test, y_pred_rf))
print("Random Forest ROC AUC Score:", roc_auc_score(y_test, y_pred_prob_rf))


# Plot ROC curve
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_prob_rf)
plt.figure(figsize=(10, 5))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (AUC = %0.2f)' % roc_auc_score(y_test, y_
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

Random Forest Accuracy: 0.9400299850074962

Random Forest Classification Report:

	precision	recall	f1-score	support
False	0.93	1.00	0.97	566
True	1.00	0.60	0.75	101
accuracy			0.94	667
macro avg	0.97	0.80	0.86	667
weighted avg	0.94	0.94	0.93	667

Random Forest ROC AUC Score: 0.9224976384564252



4.3: Model 3 (Random Forest Model with HyperParameters Tuning)

In this Particular instance, Grid Search

```
In [ ]: # Define parameter grid for Random Forest
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=RandomForestClassifier(random_state=42), param_grid=param_grid)

# Fit the grid search to the data
grid_search.fit(X_train_scaled, y_train)

# Get the best model from grid search
best_rf_model = grid_search.best_estimator_

# Make predictions
y_pred_best_rf = best_rf_model.predict(X_test_scaled)
y_pred_prob_best_rf = best_rf_model.predict_proba(X_test_scaled)[:, 1]

# Evaluate the best model
print("Best Random Forest Accuracy:", accuracy_score(y_test, y_pred_best_rf))
print("Best Random Forest Classification Report:\n", classification_report(y_test, y_pred_best_rf))
print("Best Random Forest ROC AUC Score:", roc_auc_score(y_test, y_pred_prob_best_rf))

# Plot ROC curve
fpr_best_rf, tpr_best_rf, _ = roc_curve(y_test, y_pred_prob_best_rf)
plt.figure(figsize=(10, 5))
plt.plot(fpr_best_rf, tpr_best_rf, label='Best Random Forest (AUC = %0.2f)' % roc_auc_score(y_test, y_pred_prob_best_rf))
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

Fitting 3 folds for each of 108 candidates, totalling 324 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 33 tasks | elapsed: 26.1s

[Parallel(n_jobs=-1)]: Done 154 tasks | elapsed: 1.4min


[Parallel(n_jobs=-1)]: Done 324 out of 324 | elapsed: 2.9min finished

Best Random Forest Accuracy: 0.9370314842578711

Best Random Forest Classification Report:

	precision	recall	f1-score	support
False	0.94	0.99	0.96	566
True	0.93	0.63	0.75	101
accuracy			0.94	667
macro avg	0.93	0.81	0.86	667
weighted avg	0.94	0.94	0.93	667

Best Random Forest ROC AUC Score: 0.9235384669208971



4.4: Model 4 K-Nearest Neighbors Model

```
In [ ]: # Train the K-Nearest Neighbors model
Model_4 = KNeighborsClassifier(n_neighbors=5)
Model_4.fit(X_train_scaled, y_train)

# Make predictions
y_pred_knn = Model_4.predict(X_test_scaled)
y_pred_prob_knn = Model_4.predict_proba(X_test_scaled)[:, 1]

# Evaluate the model
print("K-Nearest Neighbors Accuracy:", accuracy_score(y_test, y_pred_knn))
print("K-Nearest Neighbors Classification Report:\n", classification_report(y_test, y_p
print("K-Nearest Neighbors ROC AUC Score:", roc_auc_score(y_test, y_pred_prob_knn))


# Plot ROC curve
fpr_knn, tpr_knn, _ = roc_curve(y_test, y_pred_prob_knn)
plt.figure(figsize=(10, 5))
plt.plot(fpr_knn, tpr_knn, label='K-Nearest Neighbors (AUC = %0.2f)' % roc_auc_score(y_
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

K-Nearest Neighbors Accuracy: 0.8470764617691154

K-Nearest Neighbors Classification Report:

	precision	recall	f1-score	support
False	0.85	0.99	0.92	566
True	0.44	0.04	0.07	101
accuracy			0.85	667
macro avg	0.65	0.52	0.49	667
weighted avg	0.79	0.85	0.79	667

K-Nearest Neighbors ROC AUC Score: 0.6905765664905713



5: Evaluation

5.1: Model 1(Baseline Model)

From this Logistic Regression model, we can deduce;

- **Imbalanced Dataset:** The dataset is imbalanced, with 566 non-churning and 101 churning customers. This imbalance affects the performance metrics, particularly for the minority class (churning customers).
- **High Precision but Low Recall for Churn:** The model has a reasonable precision for predicting churn (58%), meaning that when it predicts churn, it is correct more than half the time. However, it struggles with recall (21%), meaning it misses a significant number of actual churning customers.
- **Strong Performance for Non-Churn:** The model performs very well for non-churning customers, with high precision (87%) and recall (97%), indicating it is very reliable in identifying customers who will not churn.
- **Weighted Metrics:** The weighted averages suggest the model has good overall performance, but the poor recall for churn indicates a need for improvement if predicting churn accurately is critical.

5.2: Model 2 (Random Forest Model)

From this Random forest Classification model, we can deduce:

- The ROC AUC score of 0.92 indicates a very good ability to distinguish between churning and non-churning customers.
- **Improved Performance:** The Random Forest model shows significantly improved performance over the logistic regression model, particularly in precision and F1-score for the True (churning) class.

- High Precision for Churn: The precision for the True class is perfect (1.00), meaning there are no false positives. Every customer predicted to churn actually churns.
- Balanced Recall: While the recall for the True class (0.60) is not perfect, it is much better than in the logistic regression model, indicating the Random Forest model is better at identifying actual churners.
- Overall Strong Metrics: The high accuracy, precision, recall, and F1-scores, along with the ROC AUC score, indicate that the Random Forest model performs very well across all key metrics.

5.3 Model 3(Tuned Random Forest Model)

For this Tuned Random Forest Model, we can deduce that:

- The ROC AUC score of 0.92 indicates a very good ability to differentiate between churning and non-churning customers.
- High Precision and Recall for Non-Churn: The model performs exceptionally well in predicting non-churning customers, with very high precision (0.94) and recall (0.99).
- Improved Performance for Churn: The precision for predicting churn is very high (0.93), meaning the model makes very few false positive errors. However, the recall (0.63) is still somewhat limited, indicating that some churning customers are not being identified.
- Balanced Metrics: The F1-scores and weighted averages suggest the model has a good balance of precision and recall for both classes, indicating robust overall performance.

5.4 Model 4(K-Nearest Neighbors Model)

For this KNN model, we can deduce that:

- High Recall for Non-Churn: The model performs very well in predicting non-churning customers, with high precision (0.85) and very high recall (0.99).
- Poor Performance for Churn: The model struggles significantly with predicting churning customers, with very low precision (0.44) and extremely low recall (0.04). This results in a very low F1-score (0.07) for the churn class.
- Imbalance Impact: The performance metrics indicate that the model is heavily biased towards the majority class (non-churning customers) and performs poorly on the minority class (churning customers).
- Low ROC AUC: The ROC AUC score of 0.69 shows the model's limited effectiveness in distinguishing between churners and non-churners.

5.5 Model Recommendation

Putting into consideration all the performance metrics for the models above, we can conclude that our best fit Model for predicting whether a customer will churn or not is MODEL 3. It has the highest ROC AUC Score demonstrating excellent discrimination capability with the highest F1 score indicating a balanced performance for precision and recall making it the preferred Model we recommend for use by SyriaTel.

6: RECOMMENDATION

So having analyzed the dataset provided for Syriatel customers, we recommend the following:

1. **Target High-Risk Customers** Focus retention efforts on customers identified as high-risk by the predictive model
Recommendation: Implement a priority system for customer service where high-risk customers receive faster and more personalized support. This can be in the form of a dedicated hotline or a special customer service team.
2. **Optimize Service Plans** High total day charges are associated with higher churn rates. Review and optimize the pricing and features of your day-time call plans. Offer customized plans with lower day-time charges for high-usage customers. Consider promotional offers such as temporary discounts or bonus minutes.
3. **Improve Customer Service Quality** A high number of customer service calls correlate with higher churn. Maybe the customers are dissatisfied with the customer service. Address the root causes of frequent customer service calls. Conduct training sessions for customer service representatives to enhance their problem-solving skills. Implement a feedback system to identify common issues and resolve them proactively.
4. **Promote Retention Programs** Customers without international plans and voice mail plans show different churn behaviors. Promote the voice mail plan more aggressively, especially to customers without it, as it might be linked to higher retention. Offer free trials or discounts on the voice mail plan to new customers or those who have recently had issues resolved through customer service and i guarantee this will help wholly with customer retention.
5. **Enhance Customer Engagement** Regularly engage with customers to gather feedback and address their concerns before they consider churning. Implement a customer engagement program where regular check-ins are scheduled. Use email or SMS surveys to gather feedback and act on it promptly.
6. **Personalized Marketing Campaigns** Certain demographics and usage patterns are linked to higher churn. Use the insights from the predictive model to tailor marketing campaigns. Develop personalized marketing campaigns targeting specific segments identified as high-risk. Offer them incentives like loyalty points, discounts, or free upgrades.
7. **Continuous Monitoring and Adaptation** Monitor the effectiveness of the implemented retention strategies and adapt them based on real-time data. Set up a dashboard to track key performance indicators (KPIs) such as churn rate, customer satisfaction, and the effectiveness of retention campaigns. Use this data to continuously refine your strategies.